

- [23] Google. 2024. Introducing Gemini 2.0: our new AI model for the agentic era. URL: <https://blog.google/technology/google-deepmind/google-gemini-ai-update-december-2024/>.
- [24] Carolin Holtermann, Paul Röttger, and Anne Lauscher. 2025. Around the world in 24 hours: Probing llm knowledge of time and place. *Preprint*, arXiv:2506.03984. URL: <https://arxiv.org/abs/2506.03984>, arXiv:2506.03984.
- [25] Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, Gianna Lengyel, Guillaume Bour, Guillaume Lample, L  lio Renard Lavaud, Lucile Saulnier, Marie-Anne Lachaux, Pierre Stock, Sandeep Subramanian, Sophia Yang, and 7 others. 2024. Mixtral of Experts. *arXiv preprint*. ArXiv:2401.04088. URL: <http://arxiv.org/abs/2401.04088>, doi:10.48550/arXiv.2401.04088.
- [26] Rebecca Killick, Paul Fearnhead, and Idris A. Eckley. 2012. Optimal detection of change-points with a linear computational cost. *Journal of the American Statistical Association*, 107(500):1590–1598. URL: <http://dx.doi.org/10.1080/01621459.2012.737745>, doi:10.1080/01621459.2012.737745.
- [27] Kimi. 2025. Kimi K2: Open Agentic Intelligence. URL: <https://arxiv.org/pdf/2507.20534>.
- [28] Aochong Oliver Li and Tanya Goyal. 2025. Memorization vs. reasoning: Updating LLMs with new knowledge. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 25853–25874, Vienna, Austria. Association for Computational Linguistics. URL: <https://aclanthology.org/2025.findings-acl.1326/>, doi:10.18653/v1/2025.findings-acl.1326.
- [29] Opher Lieber, Barak Lenz, Hofit Bata, Gal Cohen, Jhonathan Osin, Itay Dalmedigos, Erez Safahi, Shaked Meirom, Yonatan Belinkov, Shai Shalev-Shwartz, Omri Abend, Raz Alon, Tomer Asida, Amir Bergman, Roman Glozman, Michael Gokhman, Avashalom Manevich, Nir Ratner, Noam Rozen, and 3 others. 2024. Jamba: A Hybrid Transformer-Mamba Language Model. *arXiv preprint*. ArXiv:2403.19887. URL: <http://arxiv.org/abs/2403.19887>, doi:10.48550/arXiv.2403.19887.
- [30] Meta. . Llama 3.1. URL: https://www.llama.com/docs/model-cards-and-prompt-formats/llama3_1/.
- [31] Meta. . Llama 4. URL: <https://www.llama.com/docs/model-cards-and-prompt-formats/llama4/>.
- [32] Meta. . meta-llama/llama-3.3-70b-instruct. URL: <https://huggingface.co/meta-llama/Llama-3.3-70B-Instruct>.
- [33] Meta. 2024. Introducing Llama 3.1: Our most capable models to date. URL: <https://ai.meta.com/blog/meta-llama-3-1/>.
- [34] Meta. 2025. The Llama 4 herd: The beginning of a new era of natively multimodal AI innovation. URL: <https://ai.meta.com/blog/llama-4-multimodal-intelligence/>.
- [35] MistralAI. Changelog. URL: <https://docs.mistral.ai/getting-started/changelog/>.
- [36] MistralAI. 2025. Medium is the new large. URL: <https://mistral.ai/news/mistral-medium-3>.
- [37] MistralAI. 2025. Mistral Medium 3.1. URL: <https://docs.mistral.ai/models/mistral-medium-3-1-25-08>.
- [38] Seyed Mahed Mousavi, Simone Alghisi, and Giuseppe Riccardi. 2024. DyKnow: Dynamically Verifying Time-Sensitive Factual Knowledge in LLMs. *arXiv preprint*. ArXiv:2404.08700 [cs]. URL: <http://arxiv.org/abs/2404.08700>, doi:10.48550/arXiv.2404.08700.
- [39] OpenAI. . GPT 3.5 Turbo. URL: <https://platform.openai.com/docs/models/gpt-3.5-turbo>.
- [40] OpenAI. . GPT 4o. URL: <https://platform.openai.com/docs/models/gpt-4o>.
- [41] OpenAI. . GPT-4o mini. URL: <https://platform.openai.com/docs/models/gpt-4o-mini>.
- [42] OpenAI. 2022. Introducing ChatGPT. URL: <https://openai.com/index/chatgpt/>.
- [43] OpenAI. 2024. GPT-4o mini: advancing cost-efficient intelligence. URL: <https://openai.com/index/gpt-4o-mini-advancing-cost-efficient-intelligence/>.
- [44] OpenAI. 2024. Hello GPT-4o. URL: <https://openai.com/index/hello-gpt-4o/>.
- [45] OpenAI, Sandhini Agarwal, Lama Ahmad, Jason Ai, Sam Altman, Andy Applebaum, Edwin Arbus, Rahul K. Arora, Yu Bai, Bowen Baker, Haiming Bao, Boaz Barak, Ally Bennett, Tyler Bertao, Nivedita Brett, Eugene Brevdo, Greg Brockman, Sebastien Bubeck, Che Chang, and 107 others. 2025. gpt-oss-120b & gpt-oss-20b Model Card. *arXiv preprint*. ArXiv:2508.10925. URL: <http://arxiv.org/abs/2508.10925>, doi:10.48550/arXiv.2508.10925.
- [46] Openchat. . OpenChat: Advancing Open-source Language Models with Mixed-Quality Data. URL: <https://github.com/imoneoi/openchat>.
- [47] Openchat. . openchat/openchat_3.5. URL: https://huggingface.co/openchat/openchat_3.5.

- [48] Oded Ovadia, Menachem Brief, Moshik Mishaeli, and Oren Elisha. 2024. Fine-Tuning or Retrieval? Comparing Knowledge Injection in LLMs. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 237–250, Miami, Florida, USA. Association for Computational Linguistics. URL: <https://aclanthology.org/2024.emnlp-main.15>, doi: 10.18653/v1/2024.emnlp-main.15.
- [49] Simtheory. Kimi K2. URL: <https://simtheory.ai/model-card/kimi-k2/>.
- [50] Chen Sun, Renat Aksitov, Andrey Zhmoginov, Nolan Andrew Miller, Max Vladymyrov, Ulrich Rueckert, Been Kim, and Mark Sandler. 2025. How new data permeates llm knowledge and how to dilute it. *Preprint*, arXiv:2504.09522. URL: <https://arxiv.org/abs/2504.09522>, arXiv:2504.09522.
- [51] Qingyu Tan, Hwee Tou Ng, and Lidong Bing. 2023. Towards Benchmarking and Improving the Temporal Reasoning Capability of Large Language Models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14820–14835, Toronto, Canada. Association for Computational Linguistics. URL: <https://aclanthology.org/2023.acl-long.828>, doi: 10.18653/v1/2023.acl-long.828.
- [52] Wang-Chiew Tan, Jane Dwivedi-Yu, Yuliang Li, Lambert Mathias, Marzieh Saeidi, Jing Nathan Yan, and Alon Halevy. 2023. TimelineQA: A Benchmark for Question Answering over Timelines. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 77–91, Toronto, Canada. Association for Computational Linguistics. URL: <https://aclanthology.org/2023.findings-acl.6>, doi: 10.18653/v1/2023.findings-acl.6.
- [53] xAI. Models and Pricing. URL: <https://docs.x.ai/docs/models>.
- [54] xAI. 2025. Grok 3 Beta — The Age of Reasoning Agents. URL: <https://x.ai/news/grok-3>.
- [55] xAI. 2025. Grok 4 Model Card. URL: <https://data.x.ai/2025-08-20-grok-4-model-card.pdf>.
- [56] Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, Jialin Wang, Yang Fan, Kai Dang, Bin Zhang, Xiong Wang, Yunfei Chu, and Junyang Lin. 2025. Qwen2.5-omni technical report. arXiv:arXiv:2503.20215.
- [57] Ying Xu, Xu Zhong, Antonio Jose Jimeno Yepes, and Jey Han Lau. 2020. Forget me not: Reducing catastrophic forgetting for domain adaptation in reading comprehension. In *2020 International Joint Conference on Neural Networks (IJCNN)*, pages 1–8. doi:10.1109/IJCNN48605.2020.9206891.
- [58] Eric Zhao, Pranjal Awasthi, and Nika Haghtalab. 2025. From style to facts: Mapping the boundaries of knowledge injection with finetuning. *Preprint*, arXiv:2503.05919. URL: <https://arxiv.org/abs/2503.05919>, arXiv:2503.05919.
- [59] Chenghao Zhu, Nuo Chen, Yufei Gao, Yunyi Zhang, Prayag Tiwari, and Benyou Wang. 2025. Is your LLM outdated? a deep look at temporal generalization. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 7433–7457, Albuquerque, New Mexico. Association for Computational Linguistics. URL: <https://aclanthology.org/2025.naacl-long.381/>, doi: 10.18653/v1/2025.naacl-long.381.
- [60] Zhiyuan Zhu, Yusheng Liao, Zhe Chen, Yuhao Wang, Yunfeng Guan, Yanfeng Wang, and Yu Wang. 2025. EvolveBench: A comprehensive benchmark for assessing temporal awareness in LLMs on evolving knowledge. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16173–16188, Vienna, Austria. Association for Computational Linguistics. URL: <https://aclanthology.org/2025.acl-long.788/>, doi: 10.18653/v1/2025.acl-long.788.